Methods of image grounding and LLM-based robotic manipulation

a CogRob project

Juan Lopez, s5156750
Thijs Lukkien, s3978389

Topic:

LLMs for robot instructions

1.
The idea

“ Compare two methods of image grounding against a ground-truth method, for LLM implementations via natural language instructions ”

2.
The background

LLM → import modules

LLM → rewards
ML → actions

LLM → general knowledge
LLM → world knowledge
LLM → planning
LLM → verification
ML → task learning

2024

ProgPrompt:
Generating Situated Robot Task Plans using Large Language Models

Language to Rewards
for Robotic Skill Synthesis

Scaling Up and Distilling Down:
Language-Guided Robot Skill Acquisition

RoboGPT:
an intelligent agent of making embodied long-term decisions for daily instruction tasks

2022

PaLM-E:
An Embodied Multimodal Language Model

Embodied Task Planning
with Large Language Models

VLM → Encode vision
LLM → Encode text

NN → Classification
LLM → Planning
LLM →

LLM → planning
LLM → module use
[…]

LLM → import modules

LLM → rewards
ML → actions

LLM → general knowledge
LLM → world knowledge
LLM → planning
LLM → verification
ML → task learning

ProgPrompt

Language to Rewards

Scaling Up and Distilling Down

2021

2024

PaLM-E

Embodied Task Planning

RoboGPT

VLM → Encode vision
LLM → Encode text

NN → Classification
LLM → Planning
LLM →

LLM → planning
LLM → module use
[...]

LLM → import modules

LLM → rewards
ML → actions

LLM → general knowledge
LLM → world knowledge
LLM → planning
LLM → verification
ML → task learning

ProgPrompt:
Generating Situated Robot Task Plans using Large

Language to Rewards

Scaling Up and

An Embodied Multimodal Language Model

with Large Langua
Models

telligent agent
odied long-term decisions
r daily instruction tasks

VLM → Encode vision
LLM → Encode text

NN → Classification
LLM → Planning
LLM →

LLM → planning
LLM → module use

Model Zoo

BLOOM - LLM → text

GR-CONVNET - CNN → grasps

SAM - MMTM → Segmentation maps

CLIP - MMLM → text

KOSMOS2 - MMLM → grounded text + referred bbox images

Additional Considerations

- Only pretrained models, no fine-tuning
- Compare their few-shots/zero-shots capabilities.
- Grasp success depends on
  - Segmentation quality
  - Classification performance
  - Grasp pose estimator

Default

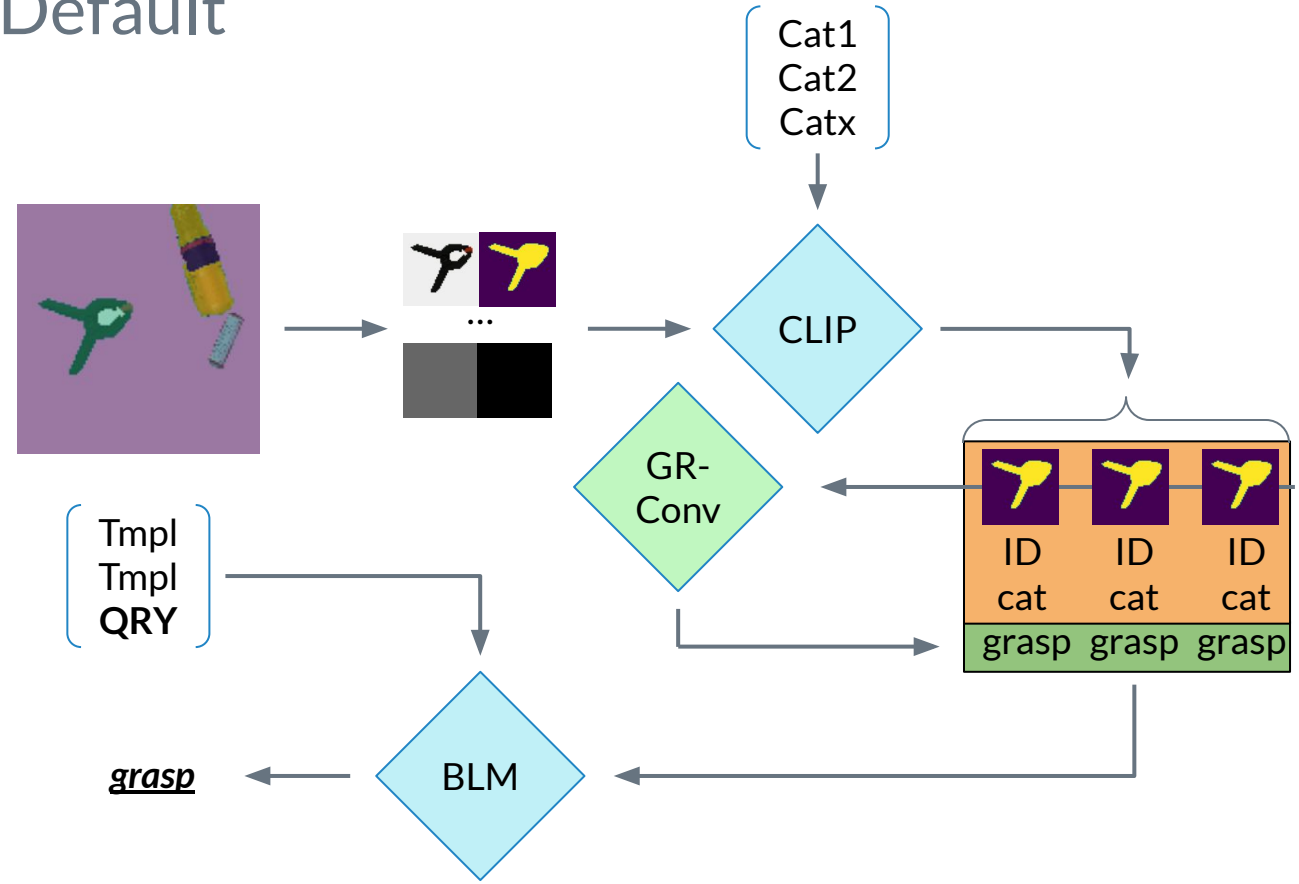
Cat1
Cat2
Catx
...
CLIP
GR-
Conv
ID
cat
ID
cat
ID
cat
grasp
grasp
grasp
Tmpl
Tmpl
QRY
BLM
grasp

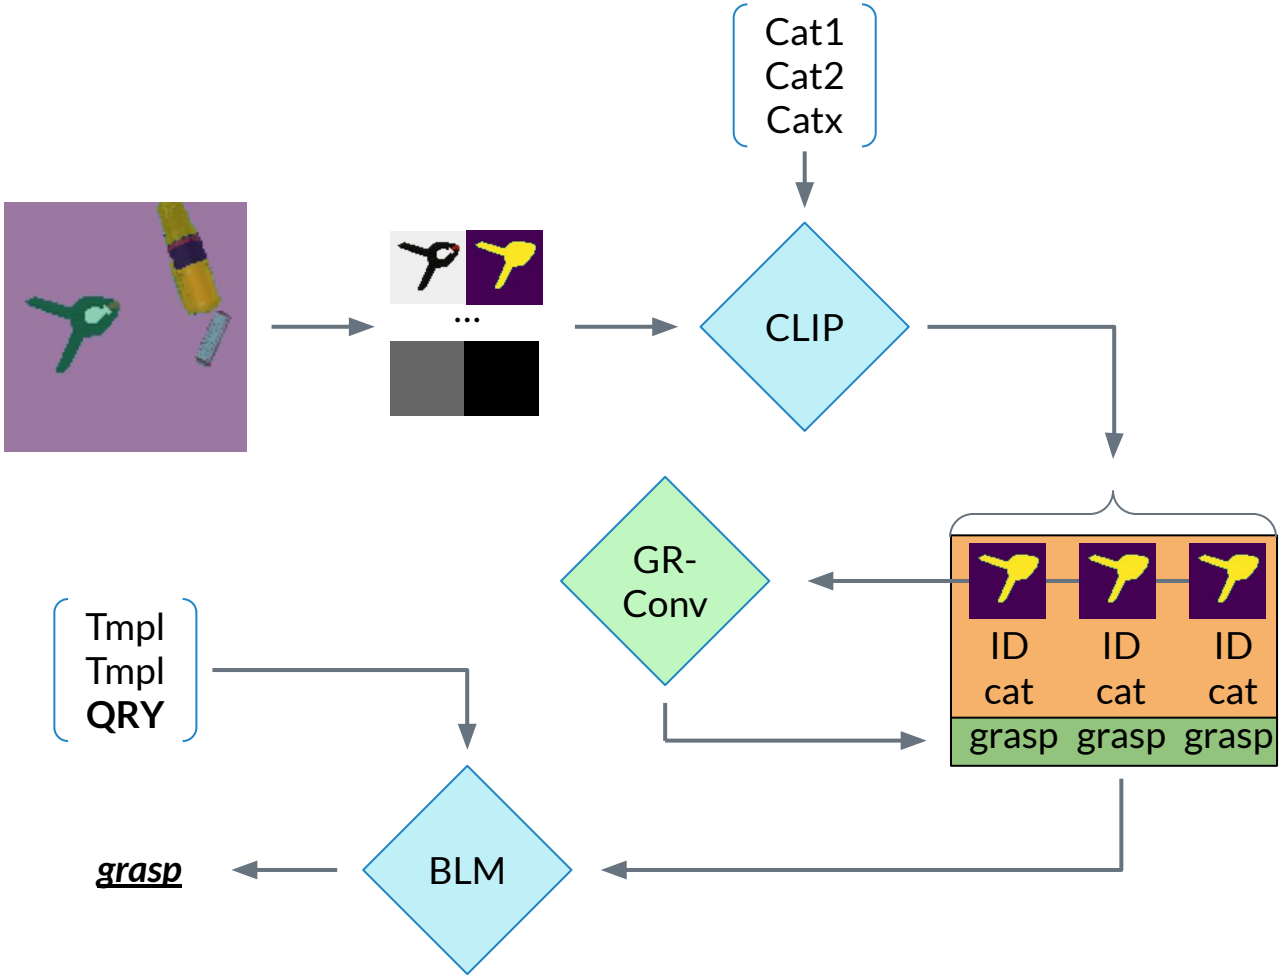
Cat1
Cat2
Catx
CLIP
GR-
Conv
ID
cat
grasp
Tmpl
Tmpl
QRY
BLM
grasp

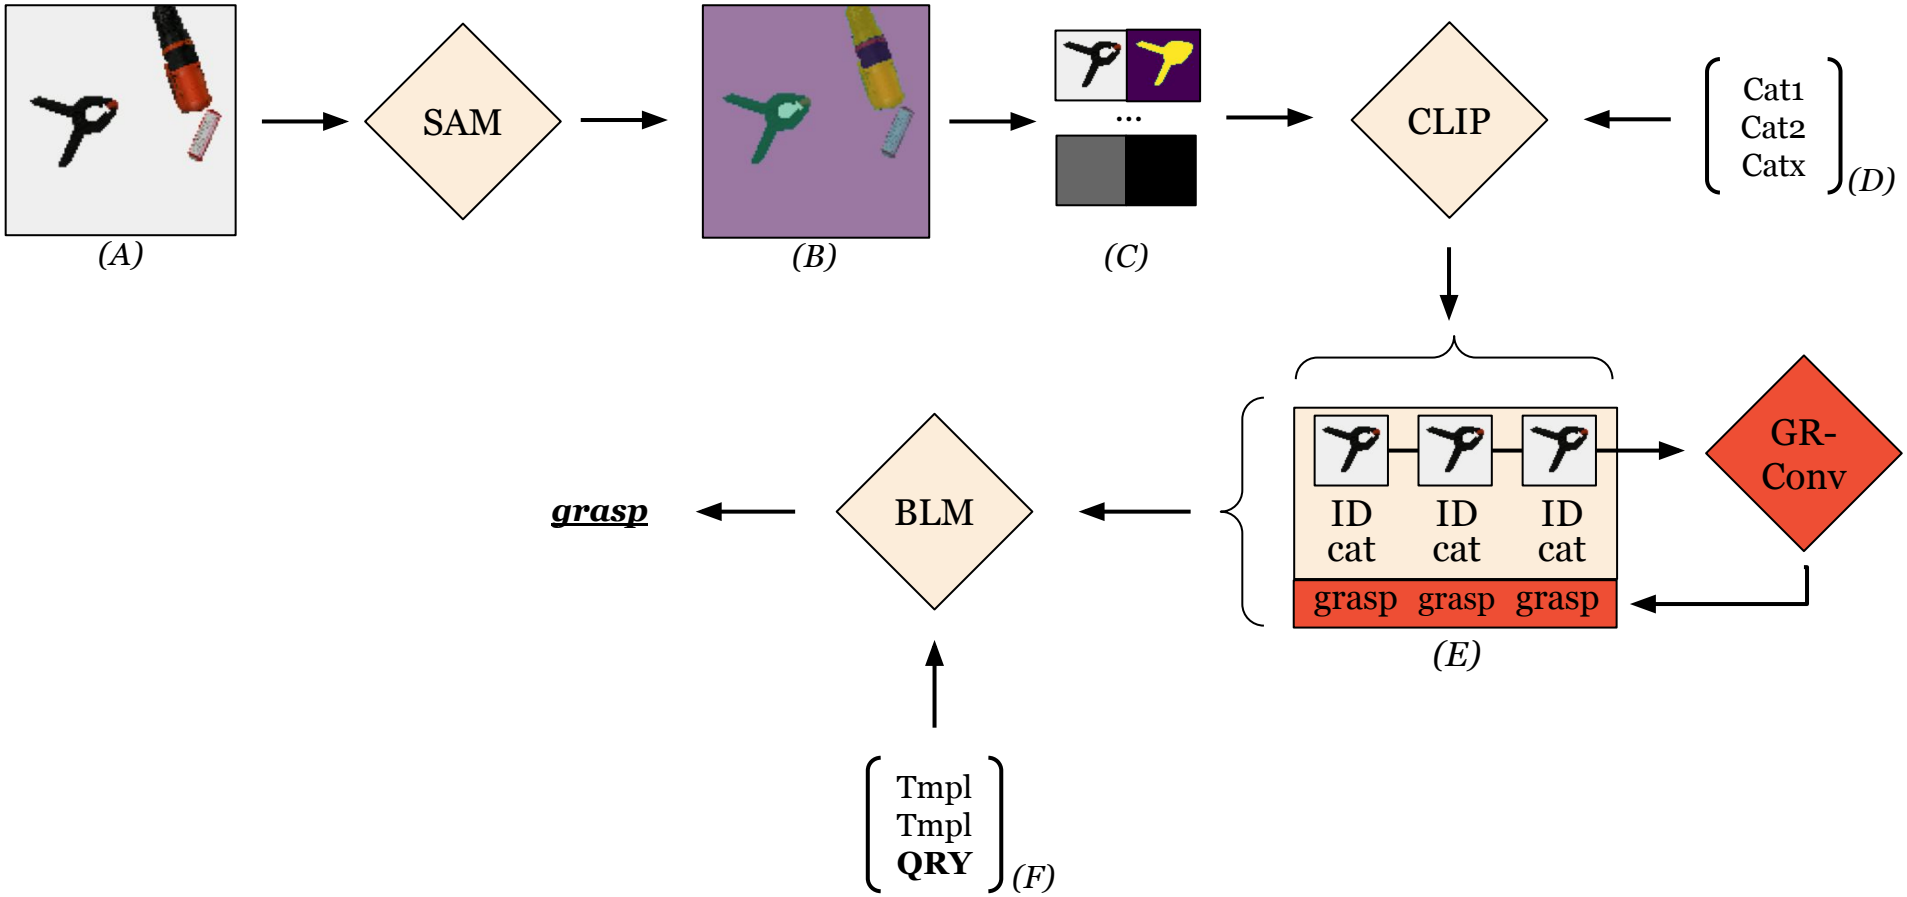
SAM
CLIP
Cat1
Cat2
Catx
(D)
(A)
(B)
(C)
ID
cat
ID
cat
ID
cat
grasp
grasp
grasp
GR-
Conv
(E)
BLM
grasp
Tmpl
Tmpl
QRY
(F)

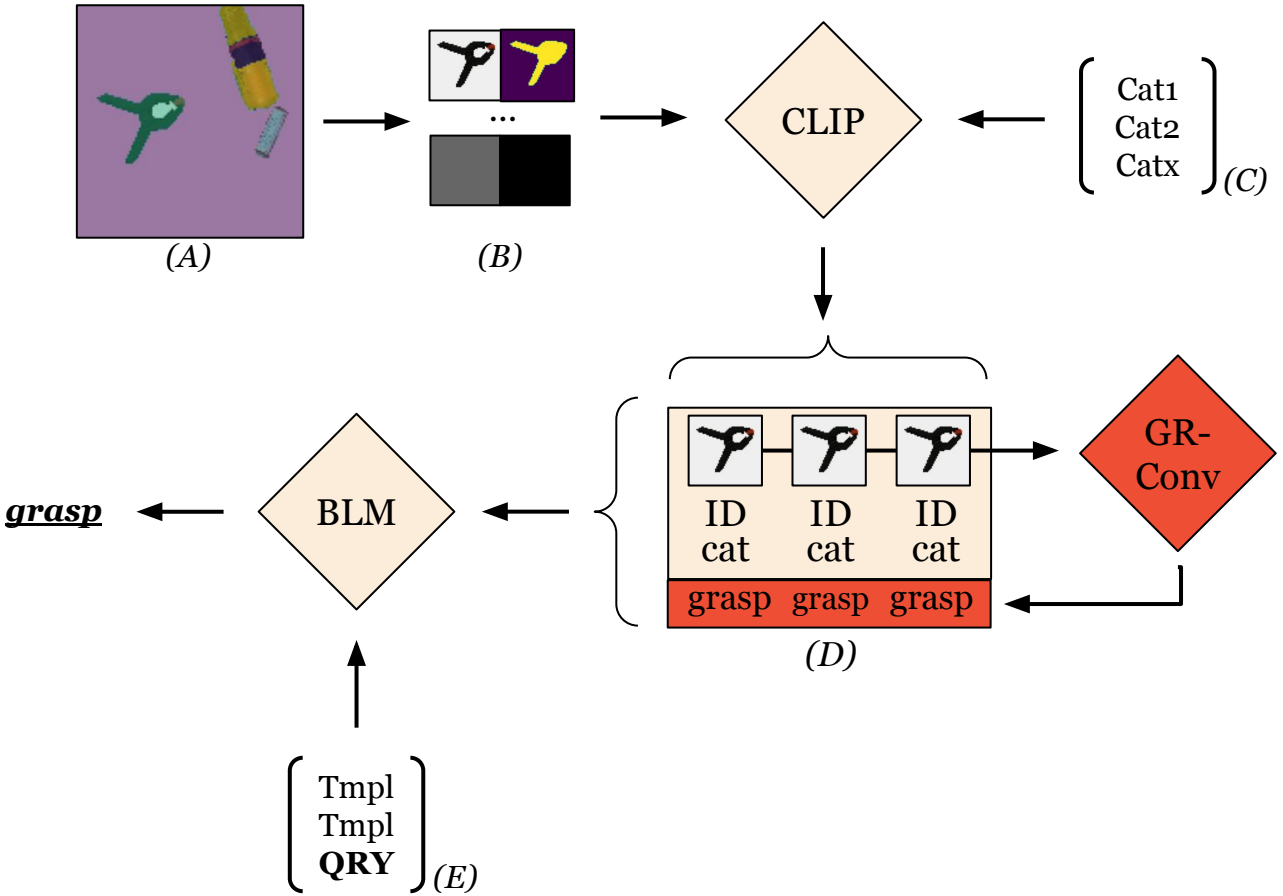
(A)
(B)
CLIP
Cat1
Cat2
Catx
(C)
ID cat
ID cat
ID cat
grasp
grasp
grasp
(D)
GR-Conv
BLM
grasp
Tmpl
Tmpl
QRY
(E)

KOSMOS-SAM

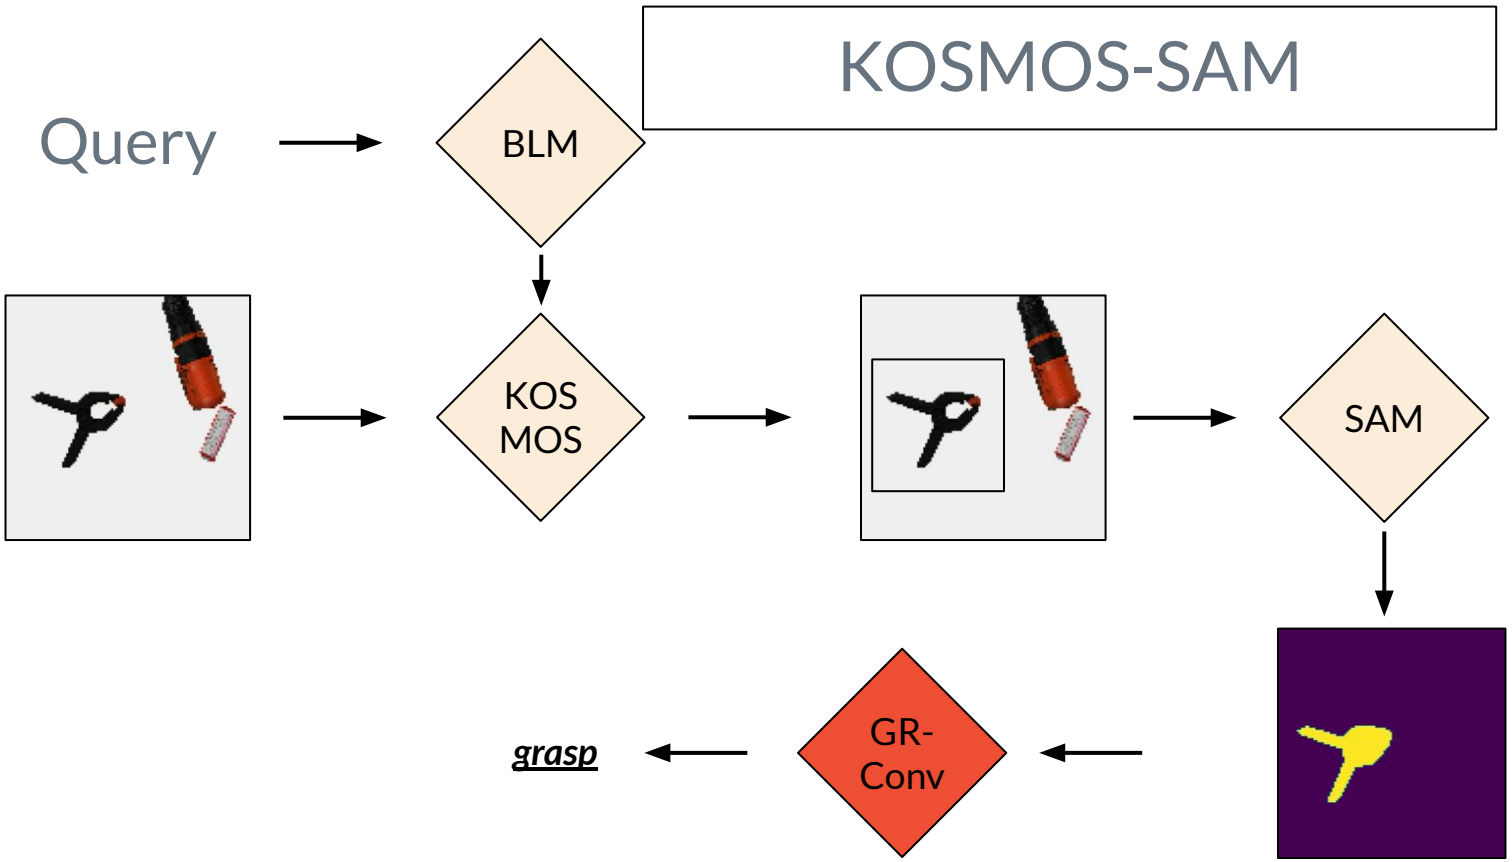

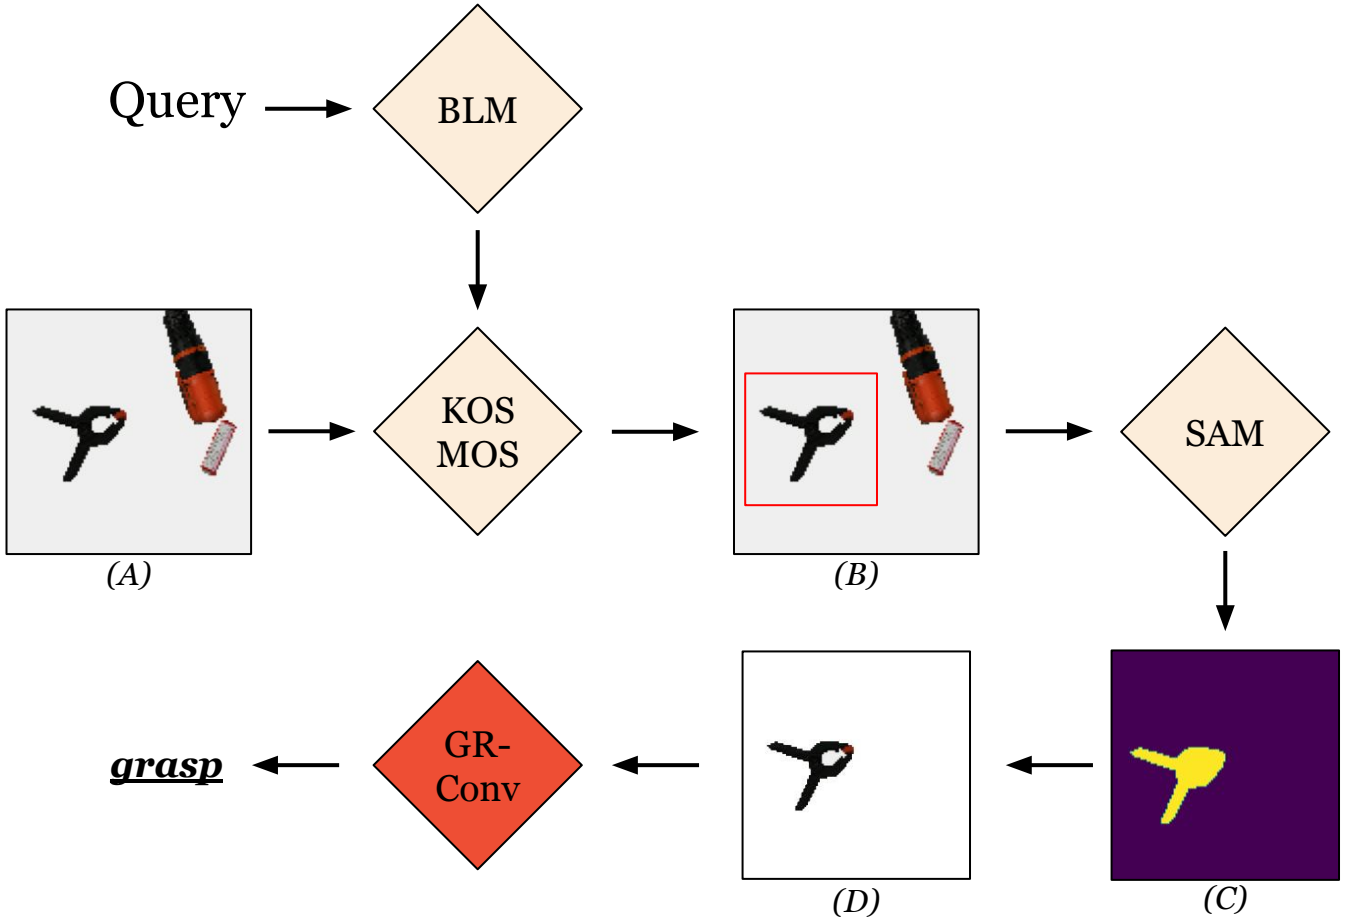
Query
BLM
KOS
MOS
(A)
(B)
SAM
(C)
(D)
GR-
Conv
grasp

3.
Evaluation

Experimental Setup

3 experimental methods:

- Default (Ground Truth Segmentation and CLIP)
- SAM (automatic) - CLIP
- KOSMOS - SAM (bounding box)

- New Collab adaptation

Experimental Setup

- 4 objects per run, 20 runs
- 1 dynamically generated query per object
- 80 queries per pipeline

What we measure

- Overall success of the query completion
- Successful grasp
- Successful delivery
- Successful Object Detection

4.
Results

Preliminary results

Kosmos-SAM > SAM-CLIP

Kosmos-SAM > ground-CLIP

CLIP-classification

ground-CLIP

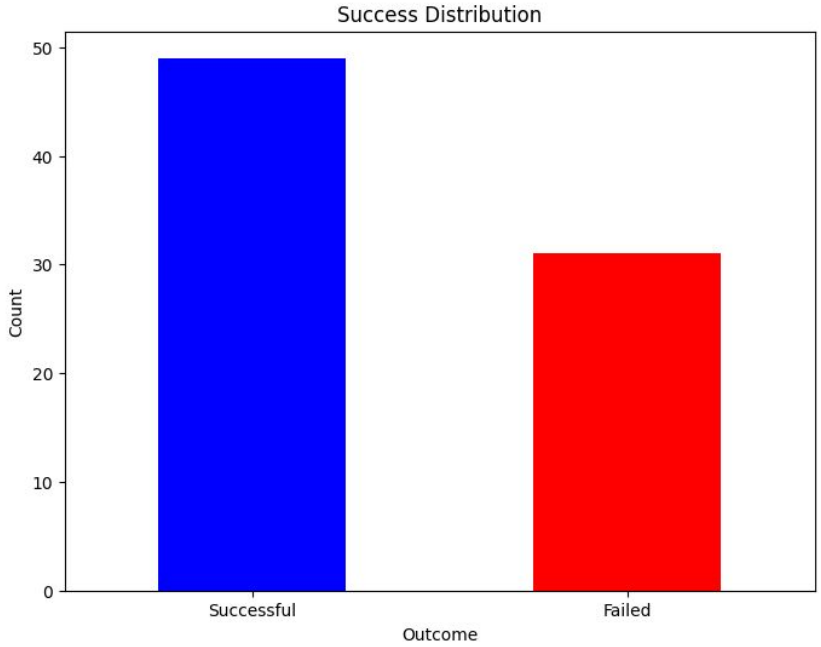
Success Distribution
Count
50
40
30
20
10
0
Successful
Failed
Outcome

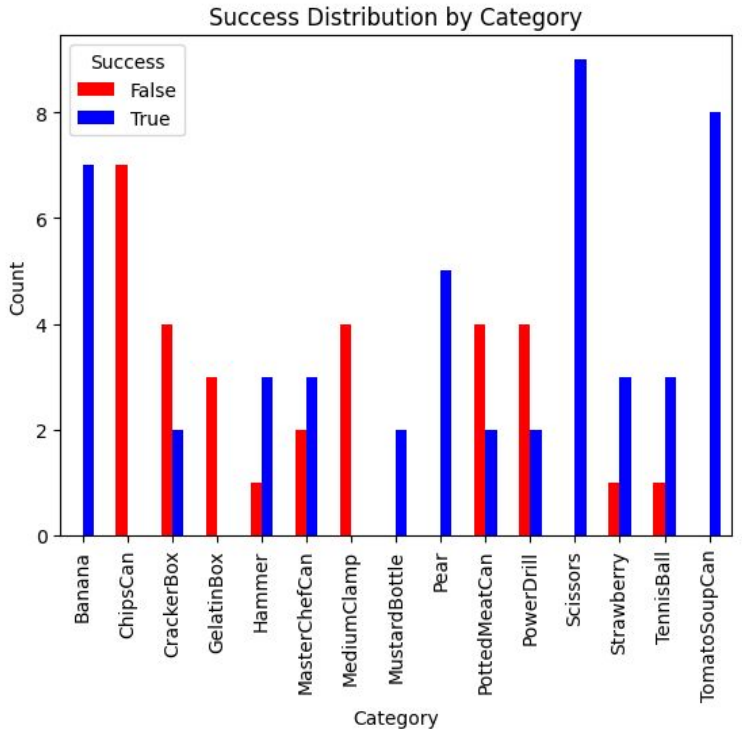
Success Distribution by Category
Success
False
True
Count
8
6
4
2
0
Banana
ChipsCan
CrackerBox
GelatinBox
Hammer
MasterChefCan
MediumClamp
MustardBottle
Pear
PottedMeatCan
PowerDrill
Scissors
Strawberry
TennisBall
TomatoSoupCan
Category

5.
Conclusion

Conclusion

Modality / model-applicability

6.
Discussion

Discussion

- CLIP and BLOOM performance (poor)
- Manual tuning
- Cascading errors : pinpointing difficulties
  - Where does low performance come from in the pipeline? Was it at the start or at the end?
- Potential for simplification
- NextChat - LLM that also creates the segmentation mask text grounding (Mixture of clip and SAM in a single model)

“Questions?”

~ Dhali ~

Extra material

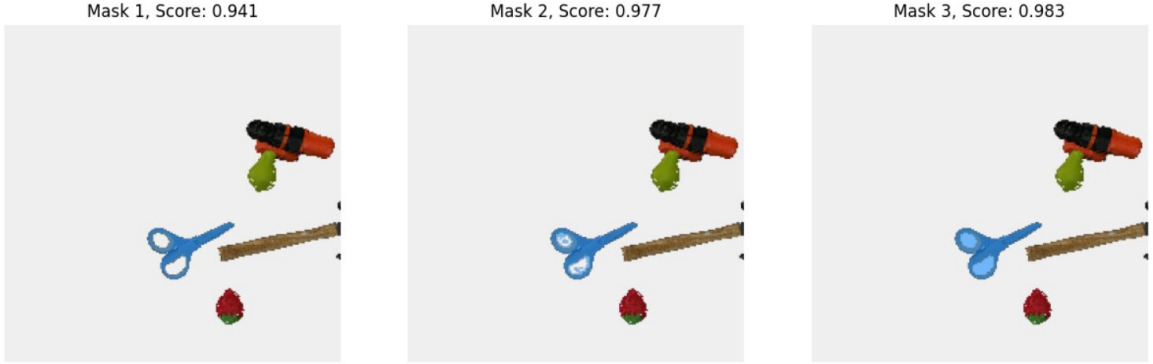


```
# put the scissors to the top left corner.
bbox = robot.ground_object("scissors")
mask = robot.segmentSAM(bbox)
robot.pick_and_place_2(mask, "top left corner")
# put the gelatin box to the top side.
bbox = robot.ground_object("gelatin box")
mask = robot.segmentSAM(bbox)
robot.pick_and_place_2(mask, "top side")
```

GR-CONVNET

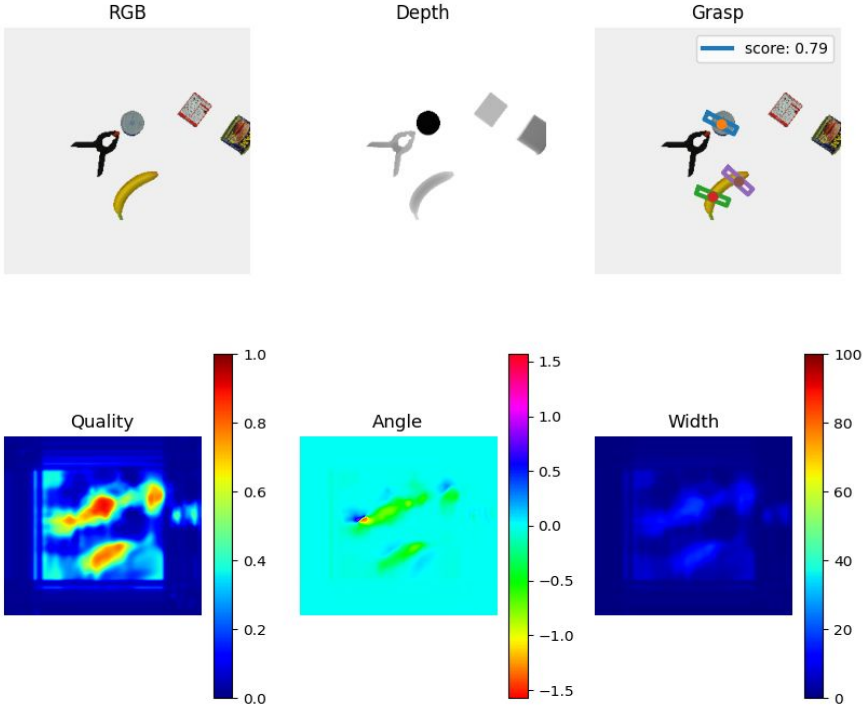
RGB
Depth
Grasp
score: 0.79
Quality
Angle
Width
1.0
0.8
0.6
0.4
0.2
0.0
1.5
1.0
0.5
0.0
−0.5
−1.0
−1.5
100
80
60
40
20
0

BLOOM

Classic LLM similar to GPT-3

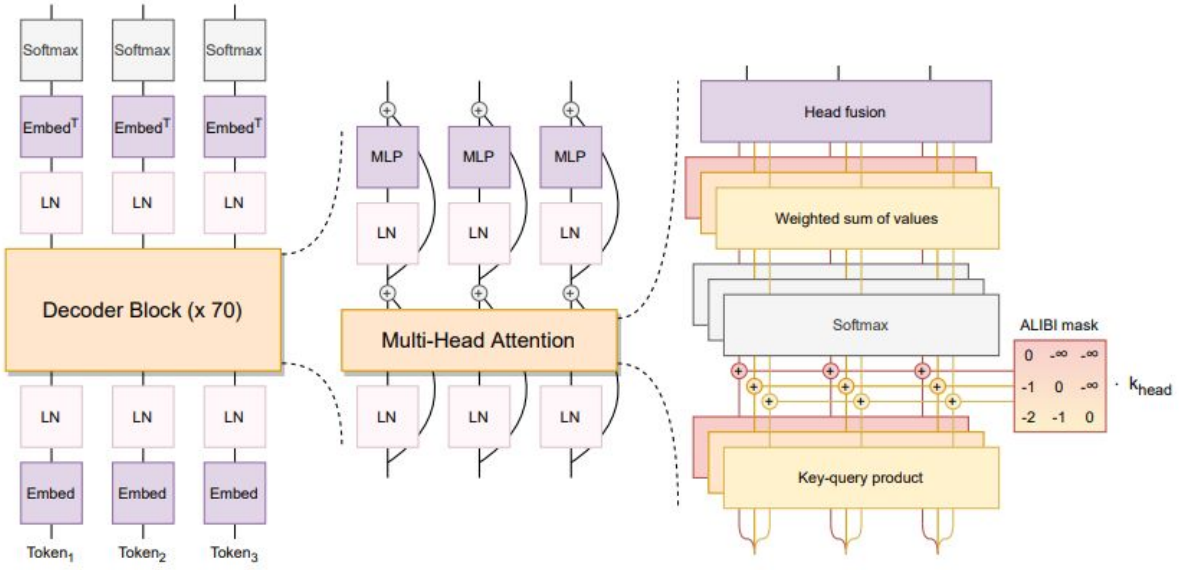

SAM

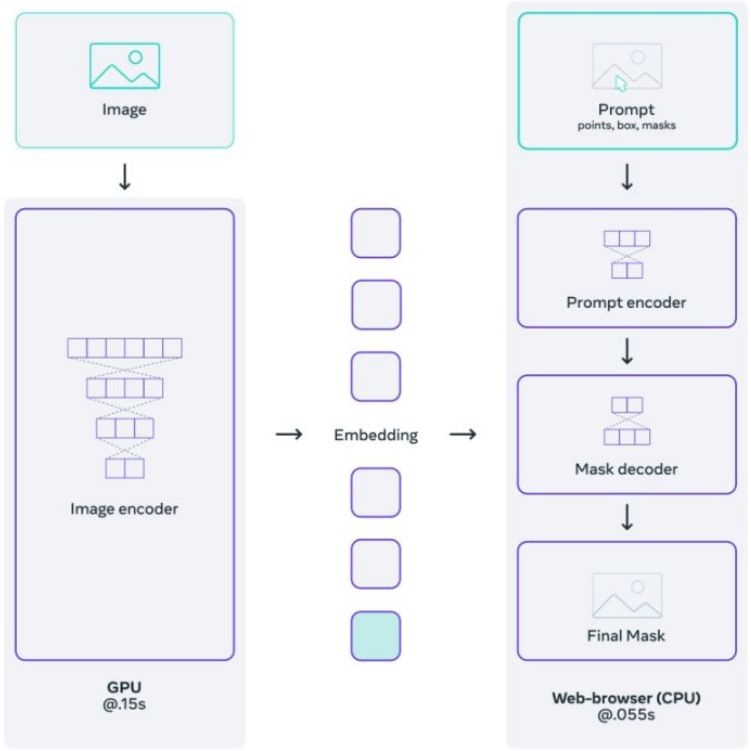
Image
Image encoder
GPU
@.15s
Embedding
Prompt
points, box, masks
Prompt encoder
Mask decoder
Final Mask
Web-browser (CPU)
@.055s

CLIP

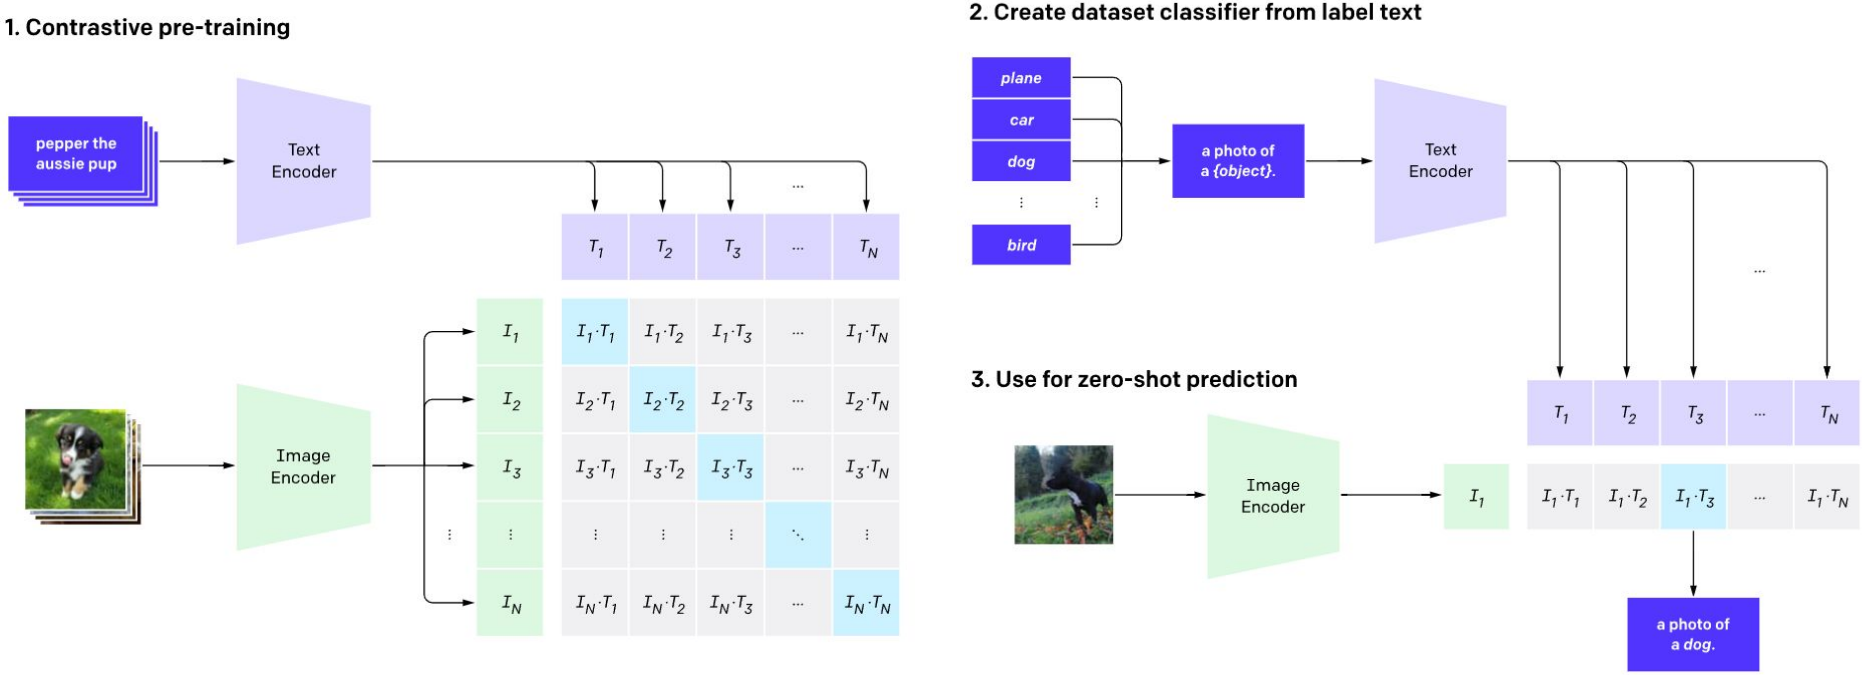
1. Contrastive pre-training
pepper the aussie pup
Text Encoder
T_1 T_2 T_3 ... T_N
Image Encoder
I_1 I_2 I_3 ⋮ I_N
$I_1 \cdot T_1$ $I_1 \cdot T_2$ $I_1 \cdot T_3$ ... $I_1 \cdot T_N$
$I_2 \cdot T_1$ $I_2 \cdot T_2$ $I_2 \cdot T_3$ ... $I_2 \cdot T_N$
$I_3 \cdot T_1$ $I_3 \cdot T_2$ $I_3 \cdot T_3$ ... $I_3 \cdot T_N$
$I_N \cdot T_1$ $I_N \cdot T_2$ $I_N \cdot T_3$ ... $I_N \cdot T_N$
2. Create dataset classifier from label text
plane
car
dog
bird
a photo of a {object}.
Text Encoder
3. Use for zero-shot prediction
Image Encoder
I_1
T_1 T_2 T_3 ... T_N
$I_1 \cdot T_1$ $I_1 \cdot T_2$ $I_1 \cdot T_3$ ... $I_1 \cdot T_N$
a photo of a dog.

Kosmos-2

Input Prompt

Question: Where is the coach?
Answer: The coach

Question: What does it say?
Answer:

Question: What makes this image weird? Answer:

Completion

The coach is standing in the middle of the field.

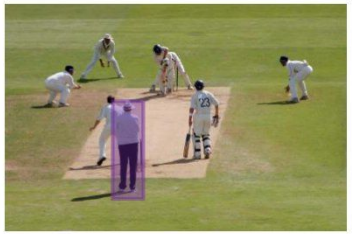

"Je suis Charlie!" is written in French on a chalkboard.

The image is weird because the snowman is sitting next to a campfire in the snow.

(1) Grounded question answering

https://colab.research.google.com/drive/1orQdJ5qkkW44uIlwPrTaKl_BfbwoiydY?usp=sharing

Once you get access to the repo